

Fitting distributional models with **brms**

Stefano Mezzini

2022-05-09

Contents

1	Setup	2
2	Fitting distributional Gaussian models with <code>mgcv</code>	4
3	Fitting distributional Gaussian models with <code>brms</code>	5
4	Fitting smooth distributional Gaussian models	9
5	Fitting smooth distributional Beta models	20

1 Setup

Start by attaching all necessary packages and creating a custom function for plotting parameters:



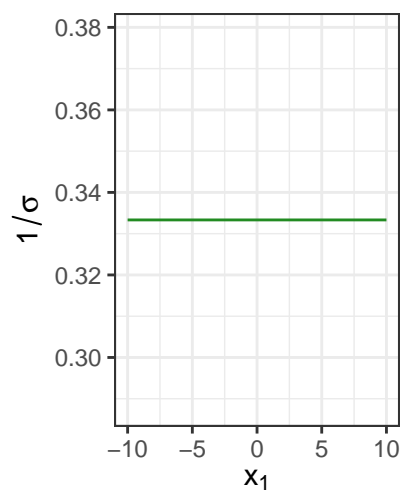




```
library('dplyr')    # for data wrangling
library('mgcv')     # for empirical Bayes modeling
library('ggplot2')  # for fancy plots
library('cowplot')  # for multi-panel fancy plots
library('gratia')   # for fancy GAM plots
library('brms')     # for Bayesian modeling
theme_set(theme_bw() + theme(legend.position = 'top')) # change default plot theme

# function to plot parameters and samples
gaus_plots <- function(.data) {
  plot_grid(
    ggplot(.data) +
      geom_point(aes(x1, y), alpha = 0.3) +
      geom_line(aes(x1, mu), color = 'darkorange') +
      labs(x = expression(x[1]), y = expression(Y ~ N(mu, ~sigma^2))),
    plot_grid(ggplot(.data) +
      geom_line(aes(x1, mu), color = 'darkorange') +
      labs(x = expression(x[1]), y = expression(mu)),
      ggplot(.data) +
      geom_line(aes(x1, sigma), color = 'dodgerblue') +
      labs(x = expression(x[1]), y = expression(sigma)),
      ggplot(.data) +
      geom_line(aes(x1, sigma2), color = 'blue') +
      labs(x = expression(x[1]), y = expression(sigma^2)),
      ggplot(.data) +
      geom_line(aes(x1, 1 / sigma), color = 'forestgreen') +
      labs(x = expression(x[1]), y = expression(1/sigma)),
      ggplot(.data) +
      geom_histogram(aes(coef_mu), fill = 'darkorange', color = 'darkorange',
        alpha = 0.3, na.rm = TRUE, bins = 20) +
      labs(x = expression(Change~'in'~mu~over~x[1]), y = 'Density'),
      ggplot(.data) +
      geom_histogram(aes(coef_sigma), fill = 'dodgerblue', color = 'dodgerblue',
        alpha = 0.3, na.rm = TRUE, bins = 20) +
      labs(x = expression(Change~'in'~sigma~over~x[1]), y = 'Density'),
      ncol = 3), ncol = 1, rel_heights = 1:2)
}
```

We can start with a simple model where μ increases linearly while σ^2 is constant. In `mgcv` we estimate $\phi = 1/\sigma$ rather than σ because ϕ is more stable than σ .

```
d <-
  tibble(x1 = seq(-10, 10, by = 0.01), # predictor variable
    mu = x1 * 1.7 - 0.54, # mean
    sigma = 3, # standard deviation
    sigma2 = sigma^2, # variance
    coef_mu = c(NA, diff(mu) / diff(x1)), # changes in mu over x1
    coef_sigma = c(NA, diff(sigma) / diff(x1)), # changes in sigma over x1
    y = rnorm(n = length(x1), mean = mu, sd = sqrt(sigma2))) # samples
gaus_plots(d)
```

Warning: position_stack requires non-overlapping x intervals



2 Fitting distributional Gaussian models with mgcv

The `logb` link function in `mgcv` is a `log` link function where `b` is the minimum σ : $\eta = \log(\sigma - b)$ and $\sigma = b + \exp(\eta)$. See `?gaulss` for more info.

```
# fit model with mgcv
m_mgcv <- gam(list(y ~ x1,
                  ~ 1),
              family = gaulss(link = list('identity', 'logb'), b = 0.01),
              data = d)
knitr::kable(summary(m_mgcv)$p.table, format = 'pipe', digits = 2)
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.54	0.07	-7.97	0
x1	1.71	0.01	144.58	0
(Intercept).1	1.11	0.02	70.27	0

The `mgcv` model has the following estimates: $\beta_0 = -0.54$, $\beta_1 = 1.71$, $\eta_\phi = 1.11 \implies \phi = b + \exp(\eta_\phi) = 3.04$.


3 Fitting distributional Gaussian models with brms

```
# check what priors can be defined by printing default ones
get_prior(formula = bf(y ~ x1,
                      sigma ~ 1, # constant sd
                      alpha ~ 0, # no skew
                      family = skew_normal(link = 'identity', link_sigma = 'log')),
          data = d)
```

```
##           prior      class coef group resp  dpar nlpar lb ub
##           (flat)         b
##           (flat)         b  x1
## student_t(3, -0.6, 12.8) Intercept
## student_t(3, 0, 2.5) Intercept          sigma
## source
## default
## (vectorized)
## default
## default
```

```
priors <-
c(
  # these priors are excessively informative to reduce fitting time
  # priors are on the link scale, so lb for sigma isn't necessary (or appropriate)
  prior(student_t(8, -0.5, 0.1), class = 'Intercept'), # student_t(df, mu, scale)
  prior(student_t(8, 1.7, 0.1), class = 'b', coef = 'x1'),
  prior(student_t(8, 1, 0.2), class = 'Intercept', dpar = 'sigma')
)
```

```
# plot priors
x <- seq(-3, 3, by = 1e-3)
layout(t(1:3))
plot(x * 0.1 - 0.5, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta[0]),
     xlab = expression(beta[0]), col = 'darkorange', lwd = 2)
plot(x * 0.1 + 1.7, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta[1]),
     xlab = expression(beta[1]), col = 'darkorange', lwd = 2)
plot(x * 0.2 + 1, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~sigma),
     xlab = expression(sigma), col = 'dodgerblue', lwd = 2)
```



```

layout(1)

# fit the model
m_brms <- brm(bf(y ~ x1,
               sigma ~ 1,
               alpha ~ 0,
               family = skew_normal(link = 'identity', link_sigma = 'log')),
             data = d,
             prior = priors,
             chains = 4,
             warmup = 100, # number of iterations to use for warmup
             iter = 2000, # total number of iterations (for warmup + sampling)
             cores = 4)

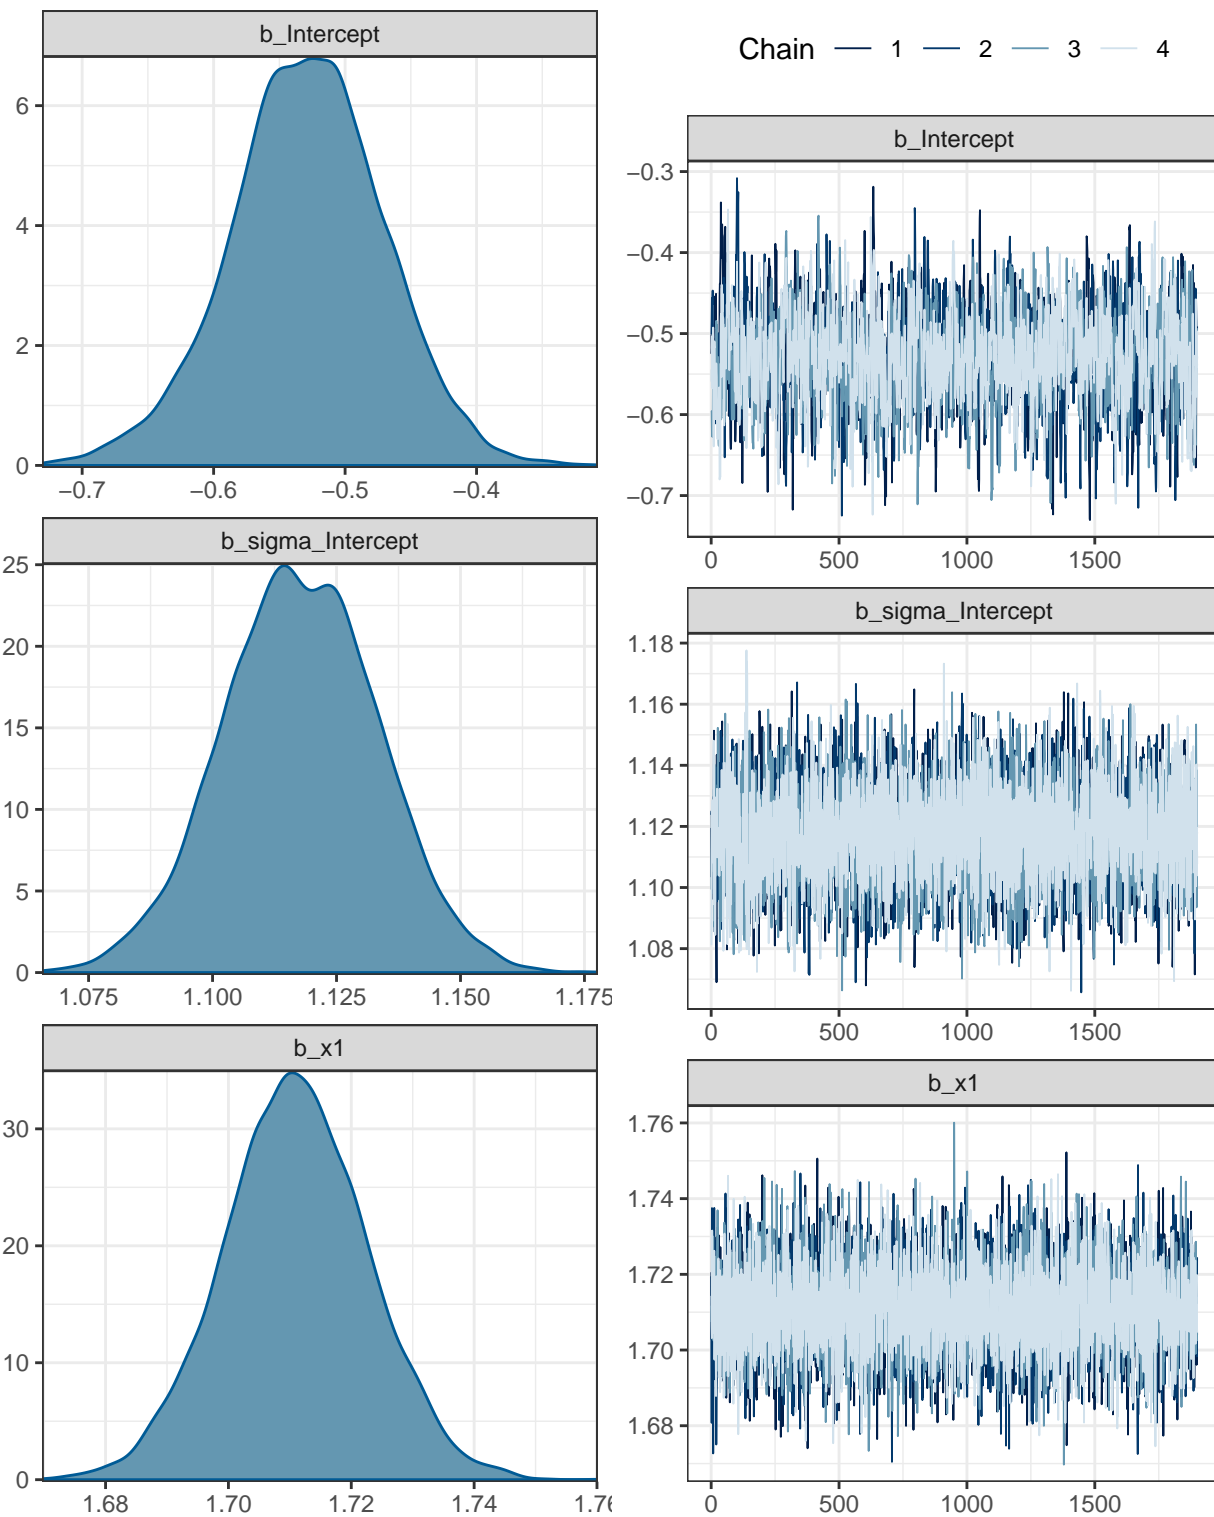
## Compiling Stan program...

## Start sampling

The posteriors are on the link scale, and  $\beta_0$  is the value of  $y$  when  $x_1 = 0$ , not when  $x_1 = \bar{x}_1$ :

plot(m_brms, N = 3) # plot posteriors for all parameters

```



$$\beta_0 = 0.54 \approx -0.529025 = \hat{\beta}_0$$


$$\beta_1 = 1.7 \approx -0.529025 = \hat{\beta}_1$$

$$\sigma = 3 \approx 3.0580636 = \hat{\sigma}$$


We can plot the parameters using `conditional_effects()`. For the trend in the mean, we can add the data

using `plot()` and specifying `points = TRUE`:


```
plot(conditional_effects(m_brms), points = TRUE)
```



```
conditional_effects(m_brms, dpar = 'sigma')
```



```
conditional_effects(m_brms, dpar = 'alpha')
```



We can extract the estimated parameters for different values of the predictor(s) using `fitted()` (or equivalently `posterior_epred()`):

```
newd <- tibble(x1 = c(-10, -5, 0, 5, 10))  
fitted(m_brms, newdata = newd) # estimated mean ( $\mu$ ) changes with  $x_1$ 
```



```
##      Estimate Est.Error      Q2.5      Q97.5
## [1,] -17.638932 0.13013998 -17.8975104 -17.3839965
## [2,]  -9.083979 0.08216094  -9.2464723  -8.9239666
## [3,]  -0.529025 0.05838712  -0.6496736  -0.4155328
## [4,]   8.025928 0.08314040   7.8612541   8.1882212
## [5,]  16.580882 0.13137818  16.3196624  16.8399042

fitted(m_brms, newdata = newd, dpar = 'sigma') # constant standard deviation (sigma)

##      Estimate Est.Error      Q2.5      Q97.5
## [1,] 3.058435 0.04765559 2.964981 3.153089
## [2,] 3.058435 0.04765559 2.964981 3.153089
## [3,] 3.058435 0.04765559 2.964981 3.153089
## [4,] 3.058435 0.04765559 2.964981 3.153089
## [5,] 3.058435 0.04765559 2.964981 3.153089

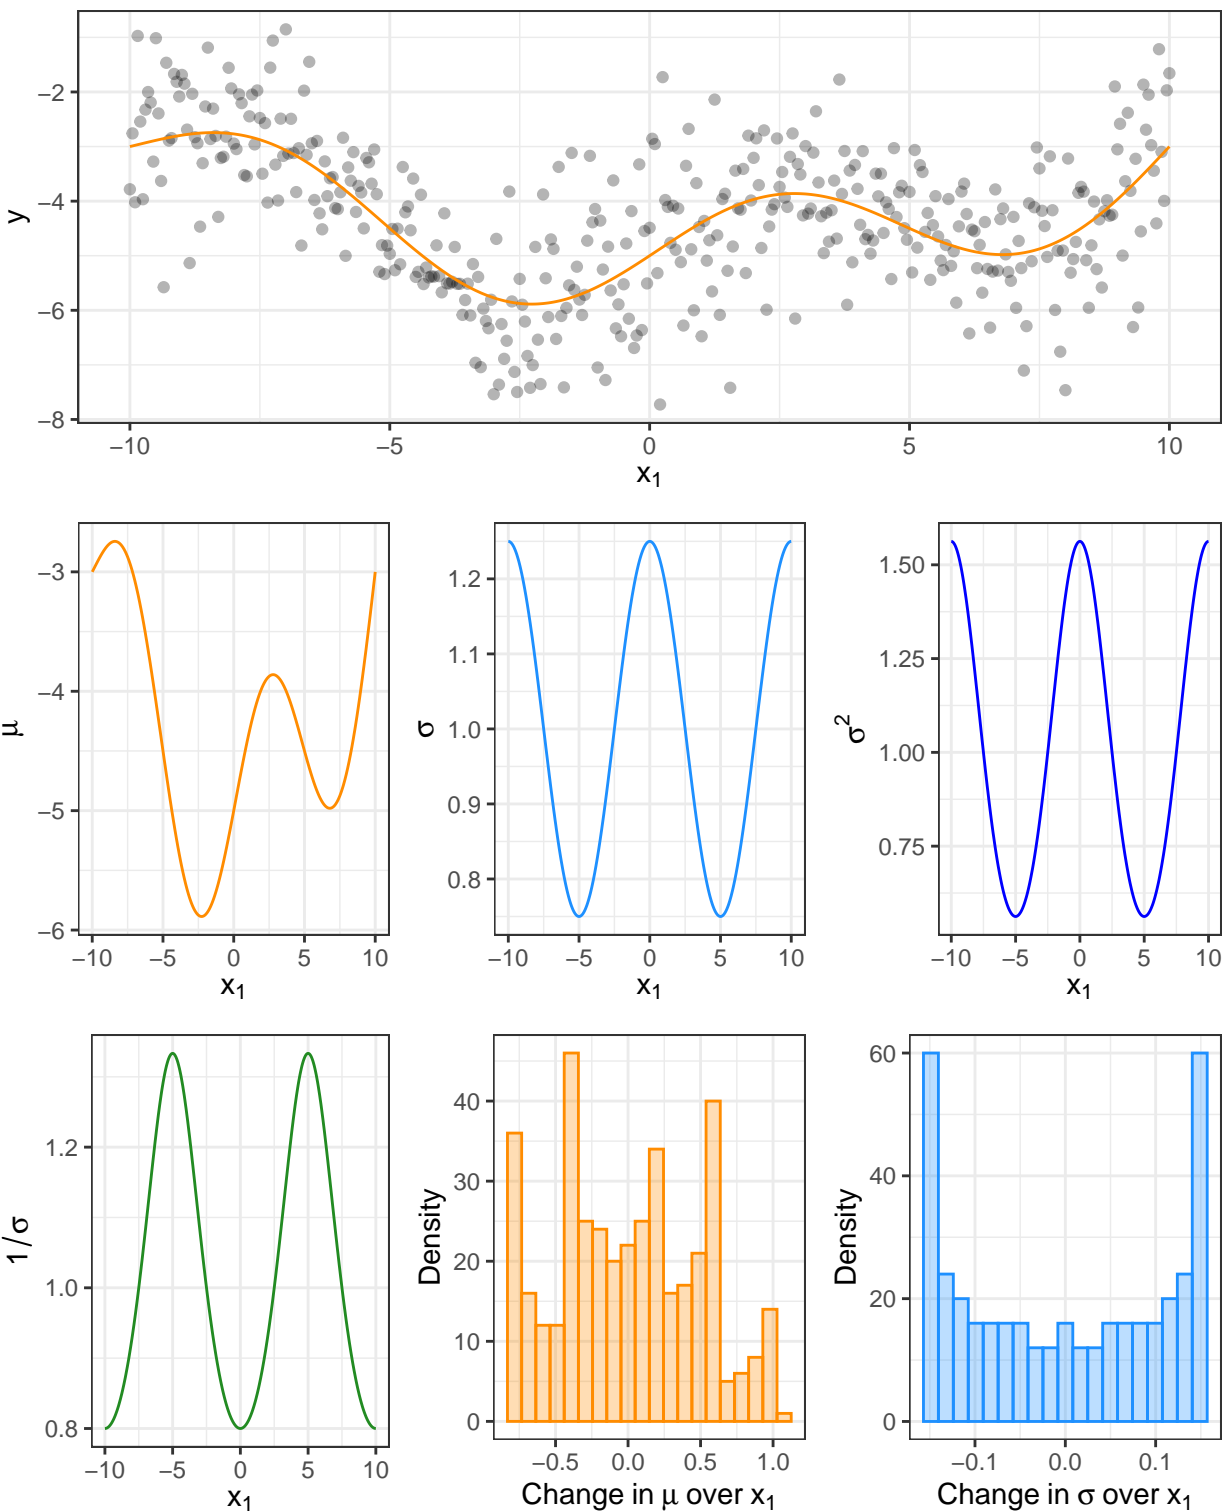
fitted(m_brms, newdata = newd, dpar = 'alpha') # skew (alpha) set to 0 (not estimated)

##      Estimate Est.Error Q2.5 Q97.5
## [1,]         0         0    0    0
## [2,]         0         0    0    0
## [3,]         0         0    0    0
## [4,]         0         0    0    0
## [5,]         0         0    0    0
```

4 Fitting smooth distributional Gaussian models

We can create a more complex “true” model where both μ and σ^2 change nonlinearly with x_1 :

```
set.seed(1) # for consistent results
d <-
  tibble(x1 = seq(-10, 10, by = 0.05), # predictor variable
         mu = sinpi(x1 / 5) + x1^2 * 0.02 - 5, # mean
         sigma = cospi(x1 / 5) * 0.25 + 1, # standard deviation
         sigma2 = sigma^2, # variance
         coef_mu = c(NA, diff(mu) / diff(x1)), # changes in mu over x1
         coef_sigma = c(NA, diff(sigma) / diff(x1)), # changes in sigma over x1
         y = rnorm(n = length(x1), mean = mu, sd = sqrt(sigma2))) # samples
gaus_plots(d)
```



We start by fitting a model with `mgcv` as before:

```
# fit model with mgcv
m_mgcv <- gam(list(y ~ s(x1, k = 20),
                  ~ s(x1, k = 15)),
              family = gaulss(link = list('identity', 'logb'), b = 0.01),
```

```


      data = d,
      method = 'REML') # restricted marginal likelihood
summary(m_mgcv)

##
## Family: gaulss
## Link function: identity logb
##
## Formula:
## y ~ s(x1, k = 20)
## ~s(x1, k = 15)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -4.29302    0.04860 -88.337  <2e-16 ***
## (Intercept).1 -0.07531    0.03600  -2.092   0.0365 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq  p-value
## s(x1)       10.289  12.55 438.75  < 2e-16 ***
## s.1(x1)      6.222   7.65  33.42 4.28e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Deviance explained =   53%
## -REML = 581.32  Scale est. = 1          n = 401

d <- bind_cols(d,
  predict(m_mgcv, type = 'link', se.fit = TRUE) %>% # mu & log(sigma-b)
  data.frame() %>%
  transmute(mu_mgcv = fit.1, # mean from mgcv model w 95% CIs
    mu_mgcv_lwr = fit.1 - 1.96 * se.fit.1,
    mu_mgcv_upr = fit.1 + 1.96 * se.fit.1,
    sigma_mgcv = exp(fit.2) + 0.01, # sigma w 95% CIs
    sigma_mgcv_lwr = exp(fit.2 - 1.96 * se.fit.2) + 0.01,
    sigma_mgcv_upr = exp(fit.2 + 1.96 * se.fit.2) + 0.01))
plot_grid(ggplot(d) +
  geom_point(aes(x1, y), alpha = 0.3) +
  geom_ribbon(aes(x1, ymin = mu_mgcv_lwr, ymax = mu_mgcv_upr, fill = 'mgcv'),
    alpha = 0.2) +
  geom_line(aes(x1, mu, color = 'true model')) +
  geom_line(aes(x1, mu_mgcv, color = 'mgcv')) +
  scale_color_brewer('Group', type = 'qual', palette = 6,
    aesthetics = c('color', 'fill')),
  ggplot(d) +
  geom_ribbon(aes(x1, ymin = sigma_mgcv_lwr, ymax = sigma_mgcv_upr,
    fill = 'mgcv'), alpha = 0.2) +
  geom_line(aes(x1, sigma, color = 'true model')) +
  geom_line(aes(x1, sigma_mgcv, color = 'mgcv')) +
  scale_color_brewer('Group', type = 'qual', palette = 6,
    aesthetics = c('color', 'fill')) +
  theme(legend.position = 'none'),

```

```
ncol = 1, rel_heights = c(1.1, 1))
```



Note that the `s.1(x1)` for σ is on the log-link scale.

We can now fit the model with `brms` and compare results. Let's start from checking the default priors:

```
# printing default priors
get_prior(formula = bf(y ~ s(x1), # smooth change in mu
                      sigma ~ s(x1), # smooth change in sigma
                      alpha ~ 0, # no skew
                      family = skew_normal(link = 'identity', link_sigma = 'log')),
          data = d)
```

```
##           prior      class  coef group resp  dpar nlpar lb ub
##           (flat)         b
##           (flat)         b sx1_1
## student_t(3, -4.2, 2.5) Intercept
## student_t(3, 0, 2.5)      sds
## student_t(3, 0, 2.5)      sds s(x1)
##           (flat)         b
##           (flat)         b sx1_1
## student_t(3, 0, 2.5) Intercept
## student_t(3, 0, 2.5)      sds
## student_t(3, 0, 2.5)      sds s(x1)
##           source
##           default
## (vectorized)
##           default
##           default
## (vectorized)
##           default
## (vectorized)
##           default
##           default
## (vectorized)
```

The population-level effects (`class = b`) for the effect of x_1 (`coef = sx1_1`) are the coefficients of the functions used to build the splines, which we can get from the `mgcv` model using `coef()`:

```
coef(m_mgcv)

## (Intercept)      s(x1).1      s(x1).2      s(x1).3      s(x1).4
## -4.29302062  0.48620408 -2.59642819 -3.56705721 -1.63468022
##      s(x1).5      s(x1).6      s(x1).7      s(x1).8      s(x1).9
##  1.19959412 -2.14919943 -0.35504812 -1.41799195 -0.84037024
##      s(x1).10     s(x1).11     s(x1).12     s(x1).13     s(x1).14
## -1.47807834  0.32902004  1.69584764 -0.50902477 -1.52757023
##      s(x1).15     s(x1).16     s(x1).17     s(x1).18     s(x1).19
## -0.64900969 -1.42082787  0.45107539 -5.73382116  3.36324822
## (Intercept).1    s.1(x1).1    s.1(x1).2    s.1(x1).3    s.1(x1).4
## -0.07531434  0.03326451 -1.64913199  0.12870659 -0.02527821
##      s.1(x1).5    s.1(x1).6    s.1(x1).7    s.1(x1).8    s.1(x1).9
## -0.16464095  0.31752956 -0.06054677 -0.42088673 -0.10861466
##      s.1(x1).10   s.1(x1).11   s.1(x1).12   s.1(x1).13   s.1(x1).14
## -0.39873199 -0.16365821  0.38895508 -1.45478851 -0.05086436
```

The coefficients for both smooths (`s(x1)` and `s.1(x1)`) range between -3 and +2.

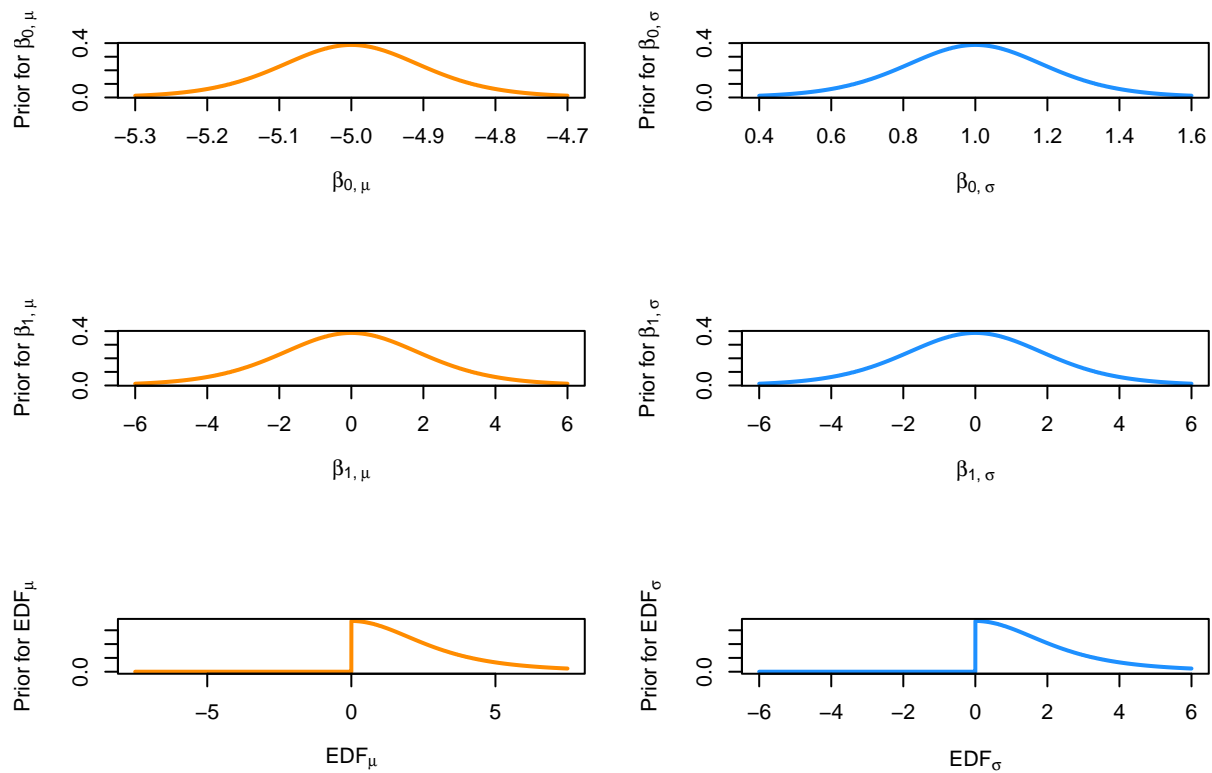
We can set priors for the intercepts as before, but now we need to also set priors for the smooth terms (instead of the linear terms). We now have coefficients for the intercept terms (β_0, μ and β_0, σ , specified with `class = 'Intercept'`), for the coefficients of the smooths (`()`):

```

priors <-
c(
  # again, priors are excessively informative to reduce fitting time
  # intercepts
  prior(student_t(8, - 5, 0.1), class = 'Intercept'),
  prior(student_t(8, 1, 0.2), class = 'Intercept', dpar = 'sigma'),
  # coefficients
  prior(student_t(8, 0, 2), class = 'b', coef = 'sx1_1'),
  prior(student_t(8, 0, 2), class = 'b', coef = 'sx1_1', dpar = 'sigma'),
  # effective degrees of freedom (a measure of wiggleness)
  prior(student_t(3, 1, 2.5), class = 'sds', lb = 0),
  prior(student_t(3, 1, 2), class = 'sds', lb = 0, dpar = 'sigma')
)

# plot priors
layout(matrix(1:6, ncol = 2, byrow = TRUE))
plot(x * 0.1 - 5, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['0','~mu]),
     xlab = expression(beta['0','~mu]), col = 'darkorange', lwd = 2)
plot(x * 0.2 + 1, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['0','~sigma]),
     xlab = expression(beta['0','~sigma]), col = 'dodgerblue', lwd = 2)
plot(x * 2, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['1','~mu]),
     xlab = expression(beta['1','~mu]), col = 'darkorange', lwd = 2)
plot(x * 2, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['1','~sigma]),
     xlab = expression(beta['1','~sigma]), col = 'dodgerblue', lwd = 2)
plot(x * 2.5, dt(x, 3) * (x > 0), type = 'l',
     ylab = expression(Prior~'for'~EDF[mu]), xlab = expression(EDF[mu]),
     col = 'darkorange', lwd = 2)
plot(x * 2, dt(x, 3) * (x > 0), type = 'l',
     ylab = expression(Prior~'for'~EDF[sigma]), xlab = expression(EDF[sigma]),
     col = 'dodgerblue', lwd = 2)

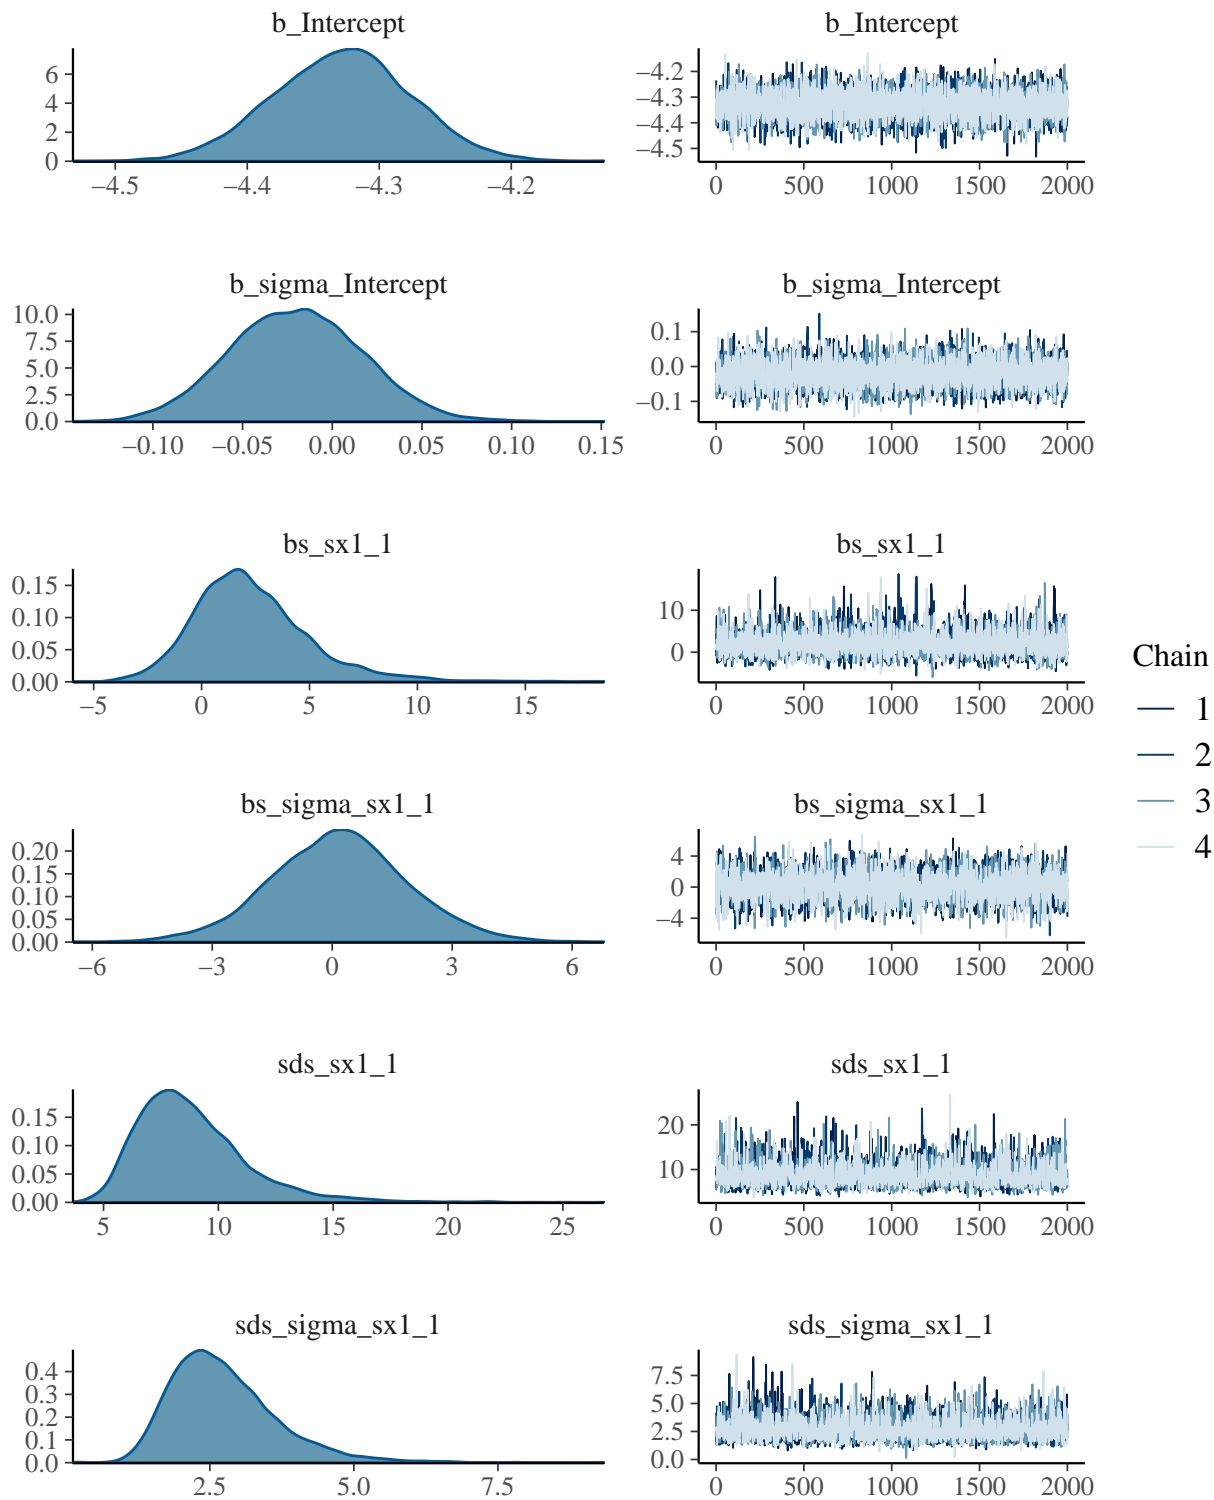
```



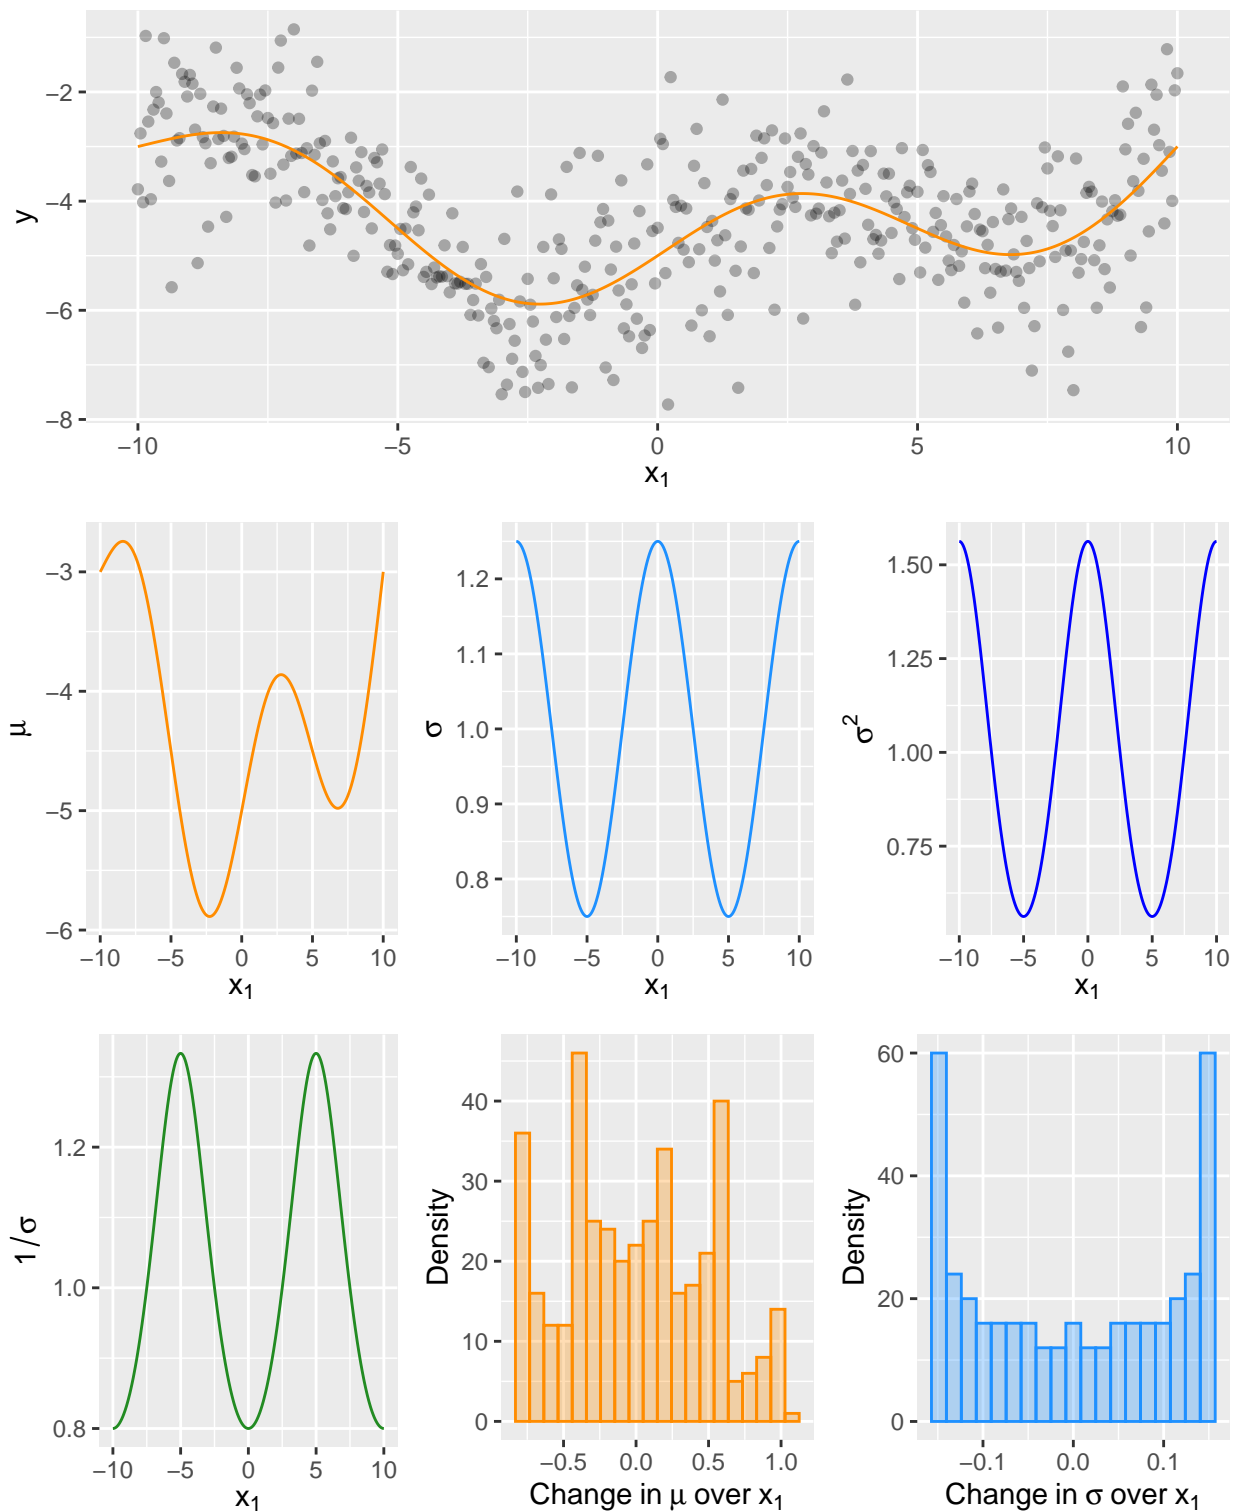
```
layout(1)

# fit the model
m_brms <- brm(bf(y ~ s(x1), # smooth change in mu
               sigma ~ s(x1), # smooth change in sigma
               alpha ~ 0, # no skew
               family = skew_normal(link = 'identity', link_sigma = 'log')),
             data = d,
             prior = priors,
             chains = 4,
             warmup = 500,
             iter = 2500,
             cores = 4,
             control = list(adapt_delta = 0.95))

## Compiling Stan program...
## Start sampling
plot(m_brms, N = 6) # plot posteriors for all parameters
```




```
gaus_plots(d) # plot the parameters to compare to the posteriors
```





We can plot the parameters using `conditional_effects()`. For the trend in the mean, we can add the data using `plot()` and specifying `points = TRUE`:


```
plot(conditional_effects(m_brms), points = TRUE)
```



```
conditional_effects(m_brms, dpar = 'sigma')
```



```
conditional_effects(m_brms, dpar = 'alpha')
```



We can extract the estimated parameters for different values of the predictor(s) using `fitted()` (or equivalently `posterior_epred()`):


```
d <- bind_cols(d,
               fitted(m_brms) %>% # estimated mean
               data.frame() %>%
```

```

      transmute(mu_brms = Estimate, # mean from brms model w 95% CIs
                mu_brms_lwr = Q2.5,
                mu_brms_upr = Q97.5),
    fitted(m_brms, dpar = 'sigma') %>% # estimated SD from brms model
    data.frame() %>%
    transmute(sigma_brms = Estimate,
              sigma_brms_lwr = Q2.5,
              sigma_brms_upr = Q97.5))

plot_grid(ggplot(d) +
  geom_point(aes(x1, y), alpha = 0.1) +
  geom_ribbon(aes(x1, ymin = mu_mgcv_lwr, ymax = mu_mgcv_upr, fill = 'mgcv'),
            alpha = 0.2) +
  geom_ribbon(aes(x1, ymin = mu_brms_lwr, ymax = mu_brms_upr,
                fill = 'brms'), alpha = 0.2) +
  geom_line(aes(x1, mu_brms, color = 'brms')) +
  geom_line(aes(x1, mu_mgcv, color = 'mgcv')) +
  geom_line(aes(x1, mu, color = 'true model')) +
  scale_color_brewer('Group', type = 'qual', palette = 6,
                    aesthetics = c('color', 'fill')),
ggplot(d) +
  geom_ribbon(aes(x1, ymin = sigma_mgcv_lwr, ymax = sigma_mgcv_upr,
                fill = 'mgcv'), alpha = 0.2) +
  geom_ribbon(aes(x1, ymin = sigma_brms_lwr, ymax = sigma_brms_upr,
                fill = 'brms'), alpha = 0.2) +
  geom_line(aes(x1, sigma_mgcv, color = 'mgcv')) +
  geom_line(aes(x1, sigma_brms, color = 'brms')) +
  geom_line(aes(x1, sigma, color = 'true model')) +
  scale_color_brewer('Group', type = 'qual', palette = 6,
                    aesthetics = c('color', 'fill')) +
  theme(legend.position = 'none'),
ncol = 1, rel_heights = c(1.1, 1))

```



5 Fitting smooth distributional Beta models

The Beta distribution has support over the interval $(0,1)$, so the mean and variance are not independent (since the a mean closer to 0 or 1 implies a smaller variance). Generally, the distribution is defined using parameters α and β , i.e., $Y \sim B(\alpha, \beta)$. With this parameterization, the mean is

$$\mu = \frac{\alpha}{\alpha + \beta},$$

and the variance is

$$\sigma^2 = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}.$$

In **brms**, the distribution is defined in terms of the mean (μ , as defined above) and the scale parameter

$$\phi = \alpha + \beta$$

. With this parameterization, the variance becomes

$$\sigma^2 = \frac{\mu(1 - \mu)}{\phi + 1}.$$

We thus can write the following functions in **R** to extract the α and β parameters from μ and ϕ :

```
get_alpha <- function(.mu, .phi) return(.mu * .phi)
get_beta  <- function(.mu, .phi) return((1 - .mu) * .phi)
```

The Beta distribution uses a logit (i.e. log of the odds) link function for μ and a log link function for ϕ . The log function and it's inverse function (the exponential) already exist in **R**, but the logit function and its inverse do not. We can then create them:

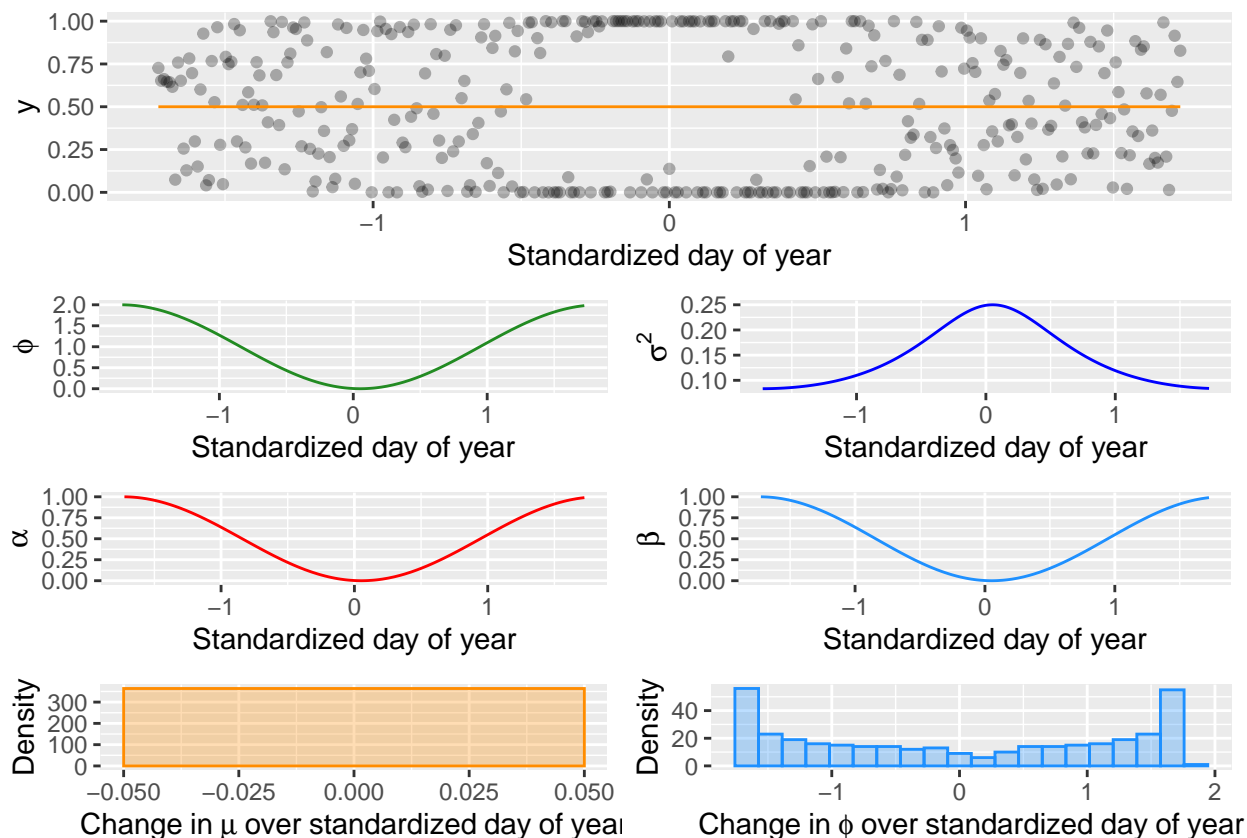
```
# logit link function and inverse function
logit <- function(.y) log(.y / (1 - .y))
inv_logit <- function(eta) exp(eta) / (exp(eta) + 1)
```

We can also create a function to create all the plots like we did for the Gaussian distribution:

```
# function to plot parameters and samples
beta_plots <- function(.data) {
  plot_grid(
    ggplot(.data) +
      geom_point(aes(doy_z, y), alpha = 0.3) +
      geom_line(aes(doy_z, mu), color = 'darkorange') +
      labs(x = 'Standardized day of year', expression(Y~'~'~B(mu,~phi))),
    plot_grid(ggplot(.data) +
      geom_line(aes(doy_z, phi), color = 'forestgreen') +
      labs(x = 'Standardized day of year', y = expression(phi)),
      ggplot(.data) +
      geom_line(aes(doy_z, sigma2), color = 'blue') +
      labs(x = 'Standardized day of year', y = expression(sigma^2)),
      ggplot(.data) +
      geom_line(aes(doy_z, alpha), color = 'red') +
      labs(x = 'Standardized day of year', y = expression(alpha)),
      ggplot(.data) +
      geom_line(aes(doy_z, beta), color = 'dodgerblue') +
      labs(x = 'Standardized day of year', y = expression(beta)),
      ggplot(.data) +
      geom_histogram(aes(coef_mu), fill = 'darkorange', color = 'darkorange',
        alpha = 0.3, na.rm = TRUE, bins = 20) +
      labs(x = expression(Change~'in'~mu~over~standardized~day~of~year),
        y = 'Density'),
      ggplot(.data) +
      geom_histogram(aes(coef_phi), fill = 'dodgerblue', color = 'dodgerblue',
        alpha = 0.3, na.rm = TRUE, bins = 20) +
      labs(x = expression(Change~'in'~phi~over~standardized~day~of~year),
        y = 'Density'),
      ncol = 2),
    ncol = 1, rel_heights = 1:2)
}
```

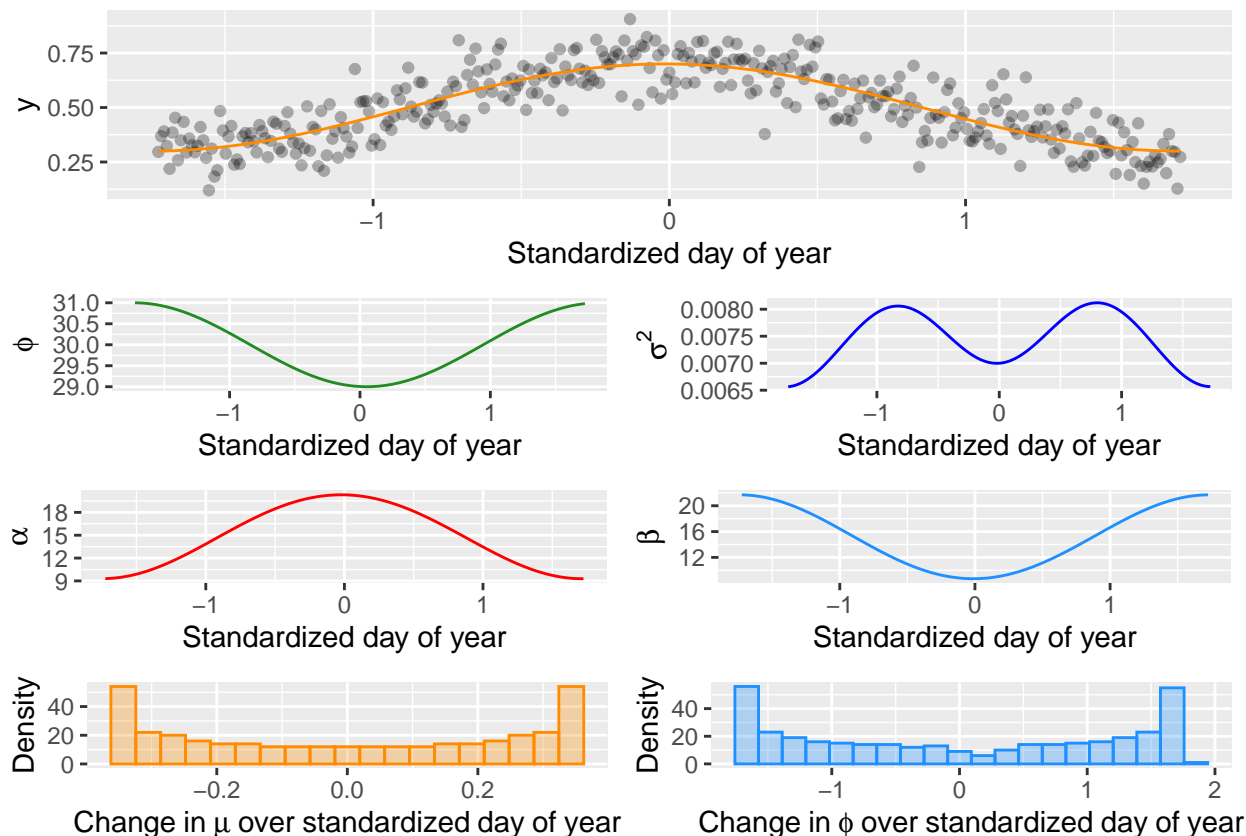
We can then create the dataset as if $Y \sim B(\alpha, \beta)$ is a linear transformation of NDVI following the linear function $Y = \frac{K+1}{2}$. We start with a simple example where μ is constant and the variance peaks on day 188.

```
tibble(doy = seq(1, 365, by = 1), # predictor variable
  doyz = (doy - mean(doy)) / sd(doy), # standardized predictor variable
  mu = 0.5, # mean; within (0, 1)
  phi = cos(doy / 60) + 1, # precision parameter; within (0, Inf)
  alpha = get_alpha(mu, phi), # shape 1
  beta = get_beta(mu, phi), # shape 2
  sigma2 = mu * (1 - mu) / (phi + 1), # variance; within (0, 0.25)
  coef_mu = c(NA, diff(mu) / diff(doyz)), # changes in mu over x1
  coef_phi = c(NA, diff(phi) / diff(doyz)), # changes in sigma over x1
  y = rbeta(n = length(doy), shape1 = alpha, shape2 = beta)) %>% # samples
mutate(y = if_else(y == 1, 0.999, y)) %>% # can't have y = 1
beta_plots()
```



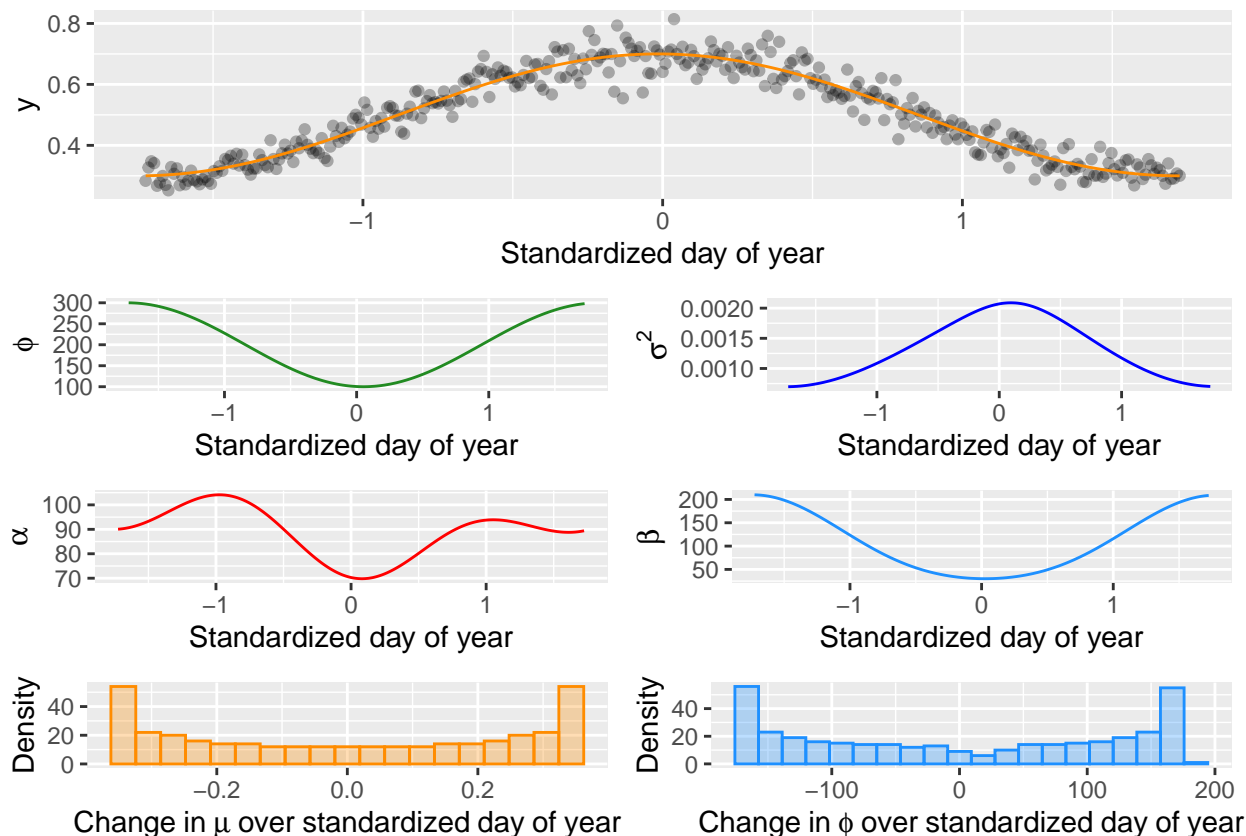
In this example, μ is highest on days 181 and 182. Note how σ^2 decreases as the mean (μ) approaches 0 or 1 because that implies σ^2 is approaching zero.

```
tibble(doy = seq(1, 365, by = 1), # predictor variable
  doy_z = (doy - mean(doy)) / sd(doy), # standardized predictor variable
  mu = sinpi((doy - 90) / 183) / 5 + 0.5, # mean; within (0, 1)
  phi = cos(doy / 60) + 30, # precision parameter; within (0, Inf)
  alpha = get_alpha(mu, phi), # shape 1
  beta = get_beta(mu, phi), # shape 2
  sigma2 = mu * (1 - mu) / (phi + 1), # variance; within (0, 0.25)
  coef_mu = c(NA, diff(mu) / diff(doy_z)), # changes in mu over x1
  coef_phi = c(NA, diff(phi) / diff(doy_z)), # changes in sigma over x1
  y = rbeta(n = length(doy), shape1 = alpha, shape2 = beta)) %>% # samples
mutate(y = if_else(y == 1, 0.999, y)) %>% # can't have y = 1
beta_plots()
```



Finally, in this example, μ peaks on day ~ 181.5 and ϕ is large and is lowest around day 188.5. We will be modeling the data from this last example. Since the days of year span a wide range of numbers, (0, 365), standardizing the predictors to Z scores will decrease the number of warm-up iterations and computation time.

```
d <-
  tibble(doy = seq(1, 365, by = 1), # predictor variable
         doy_z = (doy - mean(doy)) / sd(doy), # standardized predictor variable
         mu = sinpi((doy - 90) / 183) / 5 + 0.5, # mean; within (0, 1)
         phi = cos(doy / 60) * 100 + 200, # precision parameter; within (0, Inf)
         alpha = get_alpha(mu, phi), # shape 1
         beta = get_beta(mu, phi), # shape 2
         sigma2 = mu * (1 - mu) / (phi + 1), # variance; within (0, 0.25)
         coef_mu = c(NA, diff(mu) / diff(doy_z)), # changes in mu over x1
         coef_phi = c(NA, diff(phi) / diff(doy_z)), # changes in sigma over x1
         y = rbeta(n = length(doy), shape1 = alpha, shape2 = beta)) %>% # samples
  mutate(y = if_else(y == 1, 0.999, y)) # can't have y = 1
beta_plots(d)
```



Although `mgcv` does not support location-scale Beta models, we can fit a location Beta model to have some approximate results:

```
m_mgcv <- gam(y ~ s(doy_z, k = 10),
  family = betar(link = 'logit'),
  data = d,
  method = 'REML')

# check what priors can be defined by printing default ones
get_prior(formula = bf(y ~ s(doy_z),
  phi ~ s(doy_z),
  family = Beta(link = 'logit', link_phi = 'log')),
  data = d)
```

```
##           prior      class      coef group resp dpar nlpar lb ub
##           (flat)         b
##           (flat)         b sdoy_z_1
## student_t(3, 0, 2.5) Intercept
## student_t(3, 0, 2.5)          sds
## student_t(3, 0, 2.5)          sds s(doy_z)
##           (flat)         b
##           (flat)         b sdoy_z_1
## student_t(3, 0, 2.5) Intercept
## student_t(3, 0, 2.5)          sds
## student_t(3, 0, 2.5)          sds s(doy_z)
##           source
##           default
```



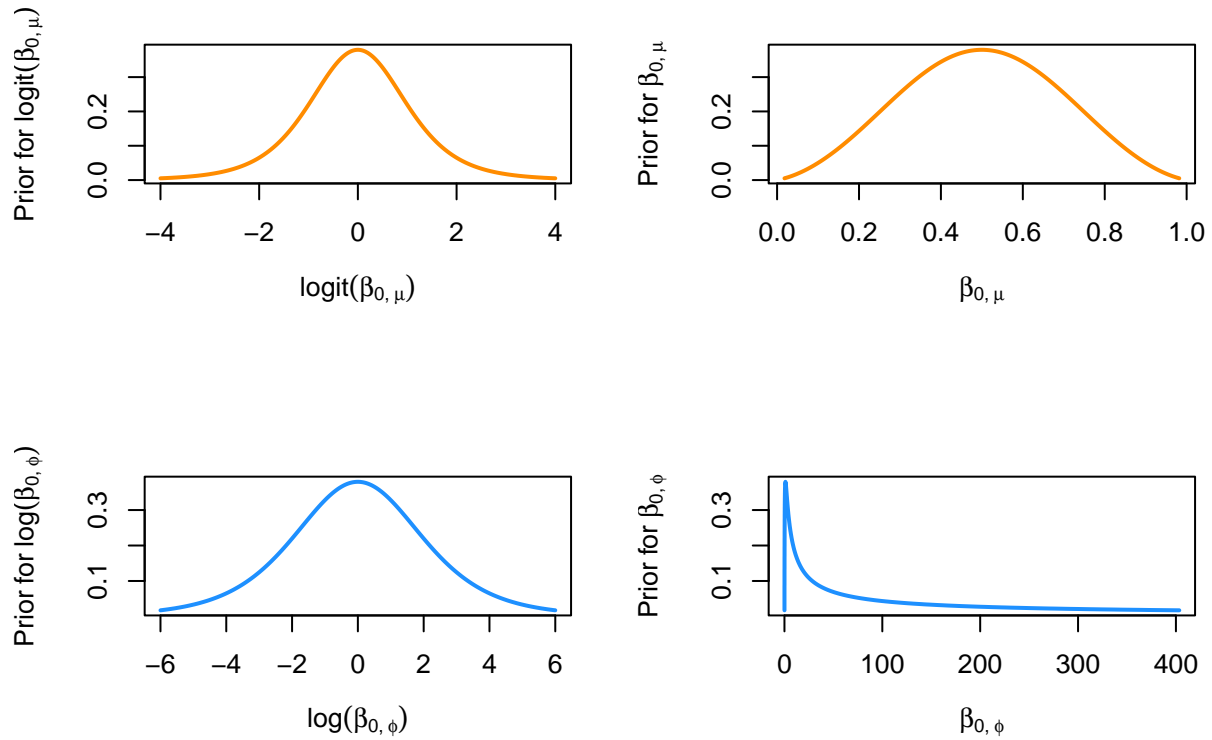
```
## (vectorized)
##      default
##      default
## (vectorized)
##      default
## (vectorized)
##      default
##      default
## (vectorized)
```

Based on “Introduction to Bayesian Statistics” by W.M. Bolstad & J.M. Curran (2017), we can find the equivalent sample sizes (n_{eq}) of each prior as $n_{eq} = \alpha + \beta + 1$ for a $B(\alpha, \beta)$ prior and $n_{eq} \sim \text{Pois}(\nu)$ for a $\Gamma(r, \nu)$ prior with shape r and rate ν .

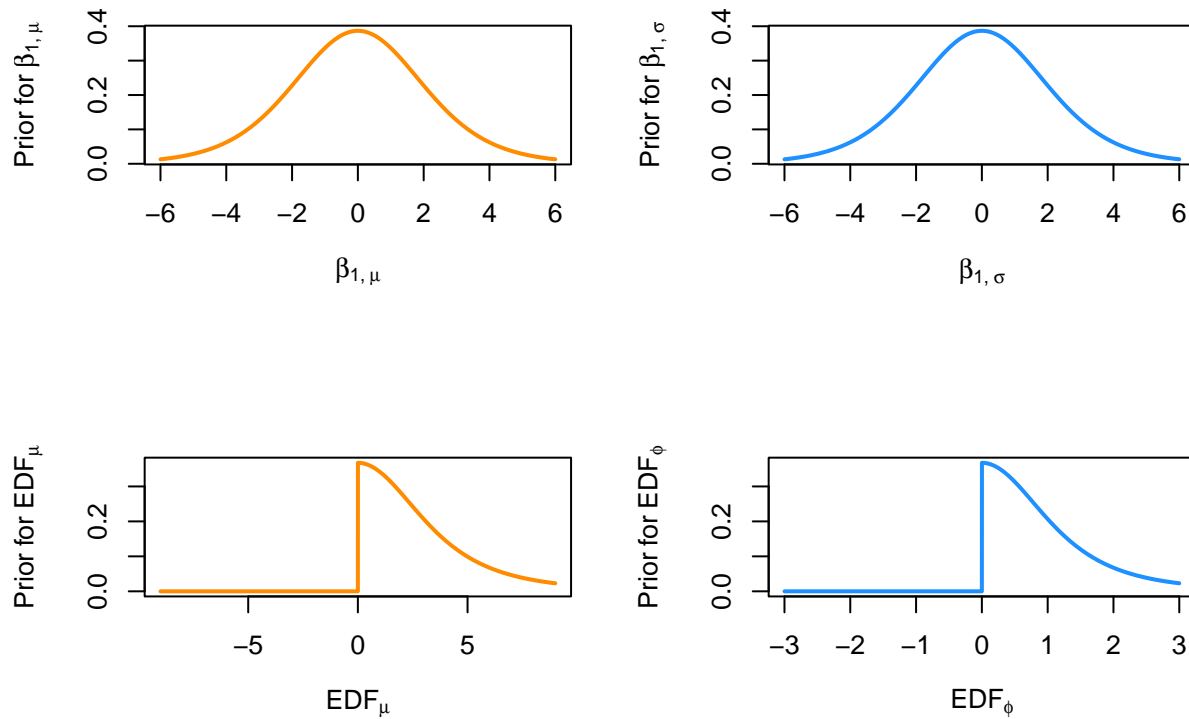
```
priors <-
c(
  # intercepts
  ## logit(mu) spans all reals, but mu is unlikely to be very near 0 or 1
  prior(student_t(5, 0, 1), class = 'Intercept'), # similar beta(3, 3) prior (n_eq = 7)
  ## log(phi) spans all reals; keep the prior for phi fairly wide
  prior(student_t(5, 0, 3), class = 'Intercept', dpar = 'phi'),
  # coefficients
  prior(student_t(8, 0, 2), class = 'b', coef = 'sdoy_z_1'),
  prior(student_t(8, 0, 2), class = 'b', coef = 'sdoy_z_1', dpar = 'phi'),
  # effective degrees of freedom (a measure of wiggleness)
  prior(student_t(3, 1, 3), class = 'sds', lb = 0),
  prior(student_t(3, 1, 1), class = 'sds', lb = 0, dpar = 'phi')
)
```

We can plot the priors to make sure they’re reasonable:

```
layout(matrix(1:4, ncol = 2, byrow = TRUE))
plot(seq(-4, 4, by = 0.01), dt(seq(-4, 4, by = 0.01), df = 5), type = 'l',
     ylab = expression(Prior~'for'~logit(beta['0,'~mu])),
     xlab = expression(logit(beta['0,'~mu])), col = 'darkorange', lwd = 2)
plot(inv_logit(seq(-4, 4, by = 0.01)), dt(seq(-4, 4, by = 0.01), 5), type = 'l',
     ylab = expression(Prior~'for'~beta['0,'~mu])),
     xlab = expression(beta['0,'~mu]), col = 'darkorange', lwd = 2)
plot(seq(-3, 3, by = 0.01) * 2, dt(seq(-3, 3, by = 0.01), df = 5), type = 'l',
     ylab = expression(Prior~'for'~log(beta['0,'~phi])),
     xlab = expression(log(beta['0,'~phi])), col = 'dodgerblue', lwd = 2)
plot(exp(seq(-3, 3, by = 0.01) * 2), dt(seq(-3, 3, by = 0.01), df = 5), type = 'l',
     ylab = expression(Prior~'for'~beta['0,'~phi]),
     xlab = expression(beta['0,'~phi]), col = 'dodgerblue', lwd = 2)
```



```
plot(x * 2, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['1','~mu]),
     xlab = expression(beta['1','~mu]), col = 'darkorange', lwd = 2)
plot(x * 2, dt(x, 8), type = 'l', ylab = expression(Prior~'for'~beta['1','~sigma]),
     xlab = expression(beta['1','~sigma]), col = 'darkorange', lwd = 2)
plot(x * 3, dt(x, 3) * (x > 0), type = 'l',
     ylab = expression(Prior~'for'~EDF[mu]), xlab = expression(EDF[mu]),
     col = 'darkorange', lwd = 2)
plot(x, dt(x, 3) * (x > 0), type = 'l',
     ylab = expression(Prior~'for'~EDF[phi]), xlab = expression(EDF[phi]),
     col = 'dodgerblue', lwd = 2)
```



```
layout(1)
```

We fit the model using `doy_z` instead of `doy` to decrease the fitting time:

```
m_brms <- brm(bf(y ~ s(doy_z),
  phi ~ s(doy_z),
  family = Beta(link = 'logit', link_phi = 'log')),
  data = d,
  prior = priors,
  chains = 4,
  warmup = 500,
  iter = 2000,
  cores = 4)
```

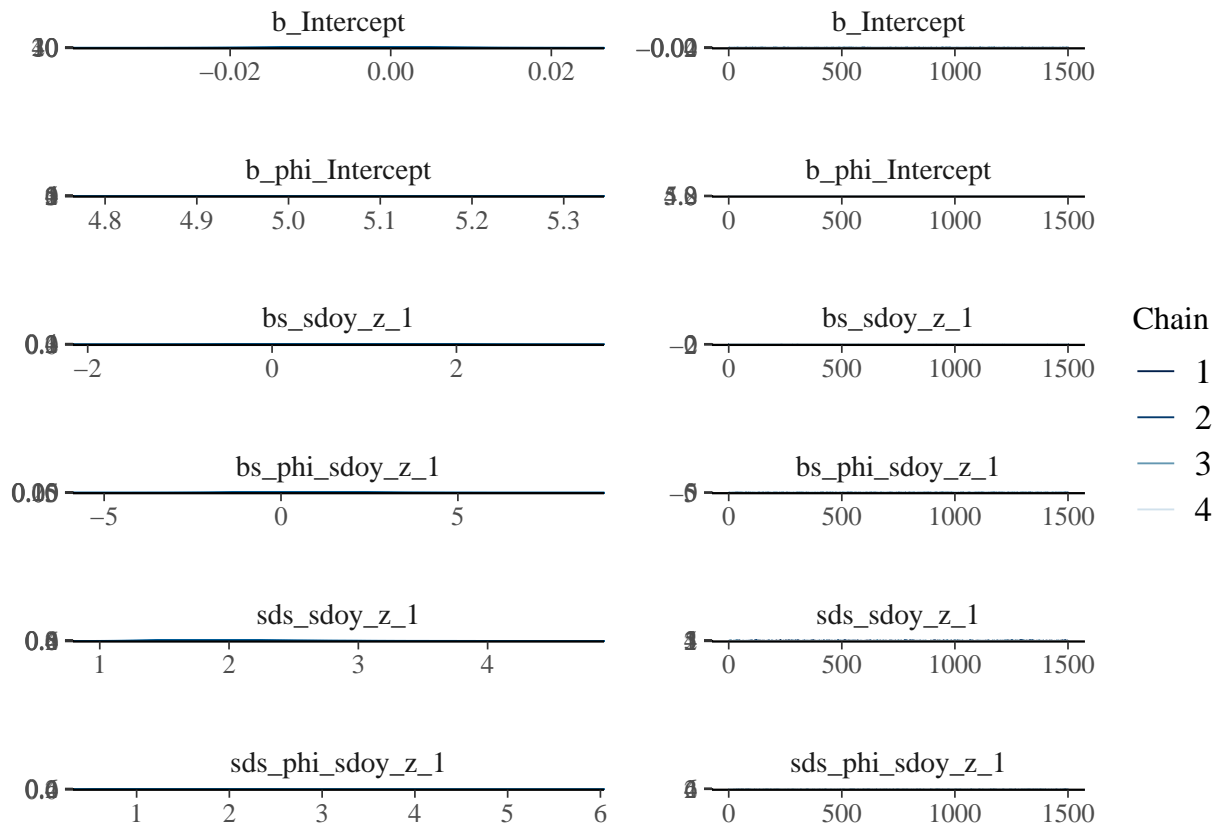
```
## Compiling Stan program...
```

```
## Start sampling
```

```
## Warning: There were 6 divergent transitions after warmup. See
## https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
## to find out why this is a problem and how to eliminate them.
```

```
## Warning: Examine the pairs() plot to diagnose sampling problems
```

```
plot(m_brms, N = 6) # plot posteriors for all 6 parameters
```



Finally, we can plot the predictions from the models:

```
d <- bind_cols(d,
  fitted(m_brms) %>% # estimated mean
  data.frame() %>%
  transmute(mu_brms = Estimate, # mean from brms model w 95% CIs
    mu_brms_lwr = Q2.5,
    mu_brms_upr = Q97.5),
  fitted(m_brms, dpar = 'phi') %>% # estimated scale from brms model
  data.frame() %>%
  transmute(phi_brms = Estimate,
    phi_brms_lwr = Q2.5,
    phi_brms_upr = Q97.5)) %>%
mutate(sigma2_brms = mu_brms * (1 - mu_brms) / (phi_brms + 1),
  sigma2_brms_lwr = mu_brms_lwr * (1 - mu_brms_lwr) / (phi_brms_lwr + 1),
  sigma2_brms_upr = mu_brms_upr * (1 - mu_brms_upr) / (phi_brms_upr + 1))

plot_grid(ggplot(d) +
  geom_point(aes(doy, y), alpha = 0.1) +
  geom_ribbon(aes(doy, ymin = mu_brms_lwr, ymax = mu_brms_upr,
    fill = 'brms'), alpha = 0.2) +
  geom_line(aes(doy, mu, color = 'true model')) +
  scale_color_brewer('Group', type = 'qual', palette = 6,
    aesthetics = c('color', 'fill')) +
```

```

  labs(x = 'Day of year', y = 'Y'),
  ggplot(d) +
    geom_ribbon(aes(doy, ymin = sigma2_brms_lwr, ymax = sigma2_brms_upr,
                    fill = 'brms'), alpha = 0.2) +
    geom_line(aes(doy, sigma2_brms, color = 'brms')) +
    geom_line(aes(doy, sigma2, color = 'true model')) +
    scale_color_brewer('Group', type = 'qual', palette = 6,
                       aesthetics = c('color', 'fill')) +
    theme(legend.position = 'none') +
    labs(x = 'Day of year', y = expression(sigma^2)),
    ncol = 1, rel_heights = c(1.1, 1))

```

